
Refinement Steps Buy Intelligibility, Not Identity: Metric-Selective Test-Time Scaling in Masked-Diffusion TTS

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 A diffusion language model buys sequential computation twice: depth, paid once
2 in parameters, and denoising steps, rented per utterance at serve time. We ask
3 whether the rented kind helps every capability equally. We train an iso-parameter
4 width–depth grid of 15 masked-diffusion codec TTS models (2 seeds, 20–125 M
5 non-embedding parameters) and sweep the refinement budget T at inference only,
6 giving a 75-point (w, d, T) surface per metric from 30 trained models. Fitting $\text{err} =$
7 $E + Aw^{-\alpha} + Bd^{-\beta} + CT^{-\tau}$ pre-registered, refinement steps attack intelligibility
8 harder than speaker identity: $\Delta\tau = \tau_{\text{WER}} - \tau_{\text{SIM}} = 0.065$, 95% run-level bootstrap
9 CI [0.053, 0.077]. Word error rate saturates only at $T^*=16$ while identity saturates
10 at 8. Depth and steps are *not* interchangeable: the substitution form $B(dT^\kappa)^{-\beta}$
11 is worse than the separable form by $\Delta\text{AICc} = +55.2$. Steps are an intelligibility
12 knob, not a general one.

1 Introduction

14 A diffusion language model can buy sequential computation twice. It *owns* depth: d transformer
15 blocks, paid once in parameters and then on every forward pass forever. It *rents* refinement steps:
16 T denoising passes over the same sequence, paid per utterance and adjustable at serve time. An
17 autoregressive model has only the first lever; its serial cost per token is fixed by the architecture the
18 moment training ends. If rented serial compute substitutes for owned serial compute, a deployment
19 can trade parameters for steps after the fact — but only for the capabilities that steps actually improve.

20 Zero-shot text-to-speech is an unusually clean probe of that question, because it separates two
21 capabilities that are measured by standard, independent metrics. Intelligibility is compositional: the
22 model must place the right phonetic content in the right order, and word error rate (WER) reads it
23 out. Speaker identity is representational: the model must carry timbre from a short prompt into new
24 content, and speaker similarity (SIM-o) reads that out. The two need not respond to the same knob.

25 We train one iso-parameter width–depth grid of masked-diffusion codec language models and then
26 sweep T at inference only, which yields a full (w, d, T) response surface per metric at zero additional
27 training cost. Two claims were pre-registered before any run. **H-D2**: refinement steps improve
28 intelligibility more than identity, i.e. test-time scaling is metric-selective. **H-D3**: for intelligibility,
29 steps and depth exchange at a measurable rate κ — how many step-doublings buy one depth-doubling.

$$\text{err}_m(w, d, T) = E_m + A_m w^{-\alpha_m} + B_m d^{-\beta_m} + C_m T^{-\tau_m} \quad \text{err}_m(w, d, T) = E_m + A_m w^{-\alpha_m} + B_m (dT^{\kappa_m})^{-\beta_m} \quad (1)$$

30 Both claims were registered before the first run, together with the thresholds, the frozen sampler
31 and the analysis plan. The answer to the first is yes and quantitative: $\Delta\tau = 0.065$ with a run-level




Figure 1: *The two ways a masked-diffusion TTS model buys sequential computation.* Phonemes and a 3-second prompt are always visible; the generated region of the $8 \times \text{frames}$ token grid starts fully masked. The backbone spends *depth* — *owned* serial compute, fixed at training time — and the sampler spends *T refinement steps per codebook level* — *rented* serial compute, chosen at serve time. One training grid over (w, d) therefore yields a whole (w, d, T) surface at inference. This paper asks whether the rented kind helps every capability equally.

bootstrap CI of $[0.053, 0.077]$, so the step exponent is reliably larger for intelligibility than for identity. The answer to the second is a clean no: the substitution form is decisively *worse* than the separable one ($\Delta\text{AICc} = +55.2$ for WER, $+89.8$ for SIM-o) and κ runs to its bound, so no finite depth-step exchange rate describes this regime. The two results together say something sharper than either alone: extra inference compute is a targeted instrument, and it cannot be treated as a substitute for architecture.

- ❶ The first per-metric (w, d, T) scaling surface for speech generation: 75 points per metric from one training grid of 30 models, released with the harness.
- ❷ $\Delta\tau = 0.065$, CI $[0.053, 0.077]$ — steps are metric-selective — and the *negative* result on depth-step substitution stated as an AICc comparison ($+55.2$) rather than a narrative, with κ reported as unidentified.
- ❸ A serving rule measured rather than assumed: c_{layer} is flat in *width* at batch 1 on the evaluation GPU, so serial latency is $8Td c_{\text{layer}}$ and the *step budget* is the only latency knob that trades against intelligibility.

2 Related work

Masked diffusion language models are now trained at scale [Nie et al., 2025, Sahoo et al., 2024, Austin et al., 2021], but their scaling is reported against a single likelihood-shaped objective, so a per-capability question cannot arise. In speech, the masked-generative recipes that made codec-token TTS practical — confidence-decoded parallel generation over residual codebook levels [Borsos et al., 2023, Chang et al., 2022] and its multi-stage successors [Wang et al., 2024] — fixed a step budget by engineering judgement and reported a single operating point; none maps error against steps as a law. Autoregressive codec TTS [Wang et al., 2023] has begun to report scaling behaviour, but in total parameters N only [Ye et al., 2025], which cannot separate width from depth and has no step axis at all; classical scaling work is likewise N -indexed [Kaplan et al., 2020, Hoffmann et al., 2022]. Test-time scaling is well studied in reasoning [Snell et al., 2024], where the payoff is measured on one correctness metric; whether extra inference compute helps different capabilities *differently* is, to our knowledge, untested in generative speech. That is the gap this abstract fills: per-metric laws in (w, d, T) for a masked-diffusion TTS model.

3 Method

Backbone and objective. A bidirectional (non-causal) pre-LN transformer over Mimi [Défossez et al., 2024] codec tokens at 12.5 Hz with the first 8 residual codebooks (2048 entries each, plus MASK and PAD). *Width* w , *depth* d , head dimension 64, GELU MLP ratio 4, RMSNorm, RoPE, no weight tying, so the non-embedding parameter count is $N = 12dw^2$. One embedding table per codebook is summed at the input and one linear head per codebook reads out. Training is SoundStorm-style [Borsos et al., 2023] coarse-to-fine masking: per example a level $l \sim \mathcal{U}\{1..8\}$ and a ratio $t \sim \mathcal{U}(0, 1]$; levels below l stay visible, level l is masked i.i.d. with probability t , levels above l are entirely MASK; the loss is cross-entropy on the masked cells of level l only. A phoneme

prefix and a 3-second prompt of codec tokens are always visible; there is no speaker encoder, so the backbone itself is the only identity pathway, and no classifier-free guidance, which would add a second inference-time axis and confound the step axis. The target length is fixed before generation from the text length and a corpus character rate (0.0602 s/char).

Frozen sampler. MaskGIT-style confidence decoding [Chang et al., 2022], level by level, with exactly T refinement steps per level, so $\text{NFE} = 8T$: a cosine schedule sets how many cells remain masked after each step, tokens are drawn categorically at temperature 1.0, and confidences are log-probabilities plus Gumbel noise annealed $1 \rightarrow 0$ within a level; committed cells stay committed. The sampler is identical for every config, seed and T — nothing is tuned per shape or per step budget. Per-item generation noise is keyed so that each item sees the same noise at the first step of every level for every T , making T the only difference between surface points of one run. Before any metric is computed we verify that T reaches the sampler: outputs at $T=1$ and $T=16$ must differ on more than 20% of generated token cells.

Design and statistics. 15 configs span three iso- N budgets $\times$ five aspect ratios, each trained at 2 seeds for 30,000 steps at effective batch 256 with μP [Yang et al., 2021] (base width 256, base LR 0.004); width transfer and depth transfer were verified by five-point LR sweeps before the grid. $T \in \{1, 2, 4, 8, 16\}$ is swept at evaluation only, giving 75 surface points per metric from 30 trained models. Metrics: WER (Whisper-large-v3 [Radford et al., 2023], greedy, the Whisper English normaliser), SIM-o (WavLM-large [Chen et al., 2022] speaker-verification embedding against the *original* prompt waveform, following Anastassiou et al. [2024]), a frozen degenerate-output rule, and UTMOS [Saeki et al., 2022] as a tertiary. Error transforms are $\text{err}_{\text{WER}} = \text{WER}$ and $\text{err}_{\text{SIM}} = 1 - \text{SIM-o}$, computed over all items including degenerates. Both forms of Eq. (1) are fitted by weighted non-linear least squares with 32 multi-starts; $\Delta\tau = \tau_{\text{WER}} - \tau_{\text{SIM}}$, κ , and $\Delta\rho = \alpha_{\text{SIM}}/\beta_{\text{SIM}} - \alpha_{\text{WER}}/\beta_{\text{WER}}$ carry percentile confidence intervals from a bootstrap that resamples at the *run* level — the five T -points of one run share a trained model, so resampling surface points would understate their correlation — with 2000 replicates recomputing every T point.

4 Results

Steps are metric-selective (H-D2, supported). Fitted on all 75 surface points, the step exponent is $\tau_{\text{WER}} = 0.809$ (CI [0.803, 0.814]) against $\tau_{\text{SIM}} = 0.744$ (CI [0.732, 0.755]), so $\Delta\tau = 0.065$ with CI [0.053, 0.077], excluding zero. The raw surface shows it without any model: over the grid, $T=1 \rightarrow T=16$ takes WER from 113.1% to 17.7% (a factor of $6.4\times$ on the error scale) while SIM-o moves only $0.208 \rightarrow 0.371$ ($1.26\times$ on $1 - \text{SIM-o}$), and the saturation steps differ: $T_{\text{WER}}^* = 16$ against $T_{\text{SIM}}^* = 8$. The power gate passed in all three budgets before these tests were run.

Depth and steps do not substitute (H-D3, not supported). The pre-registered substitution form collapses: $\Delta\text{AICc}(M_{\text{sub}} - M_{\text{sep}}) = +55.2$ for WER and $+89.8$ for SIM-o, where the hypothesis required $\leq +4$, and the fitted κ sits on its bound for both metrics, i.e. the exchange rate is not identified. We therefore quote no κ . The mechanism is visible in the fits: T enters with a large additive term that decays as $T^{-\tau}$ towards a floor, whereas depth changes the floor itself; a model that forces the two into one factor dT^κ cannot fit both the strong depth ordering at fixed T and the strong T ordering at fixed depth.

Shape matters, but the anisotropy ratio is not identified (H-D1, not supported). At $T=16$ the two-axis model beats the N -only model for both metrics ($\Delta\text{AICc} = -29.5$ for WER, -16.3 for SIM-o, against a -4 threshold), so error is not a function of parameter count alone — iso- N shape choice matters. The direction of the anisotropy, however, we cannot claim: $\Delta\rho = -0.615$ with CI [-1.005, 1.011], spanning zero, and the width coefficient A collapses to its lower bound for both metrics at $T=16$, which leaves α — and hence $\rho = \alpha/\beta$ — unidentified. What the data do show unambiguously is a strong *depth* ordering at fixed parameter count: the deepest shape beats the shallowest by 33.5, 16.5 and 5.6 WER points in the 20 M, 50 M and 125 M budgets, while SIM-o varies by under 0.065 across all 15 shapes.

Extrapolation and reliability. Fitting the two smaller budgets and predicting the largest gives MAPE 2.2% for SIM-o (versus 3.0% for the N -only model), inside the 15% the hypothesis asked




Figure 2: *Test-time scaling is metric-selective.* Error versus **refinement steps** for all 15 configs, each normalised to its own $T=1$ value; thick lines are grid means and the dotted line marks the pooled saturation step T^* . **(A)** WER is still falling at $T=16$ ($T_{\text{WER}}^* = 16$). **(B)** SIM-o has flattened by $T=8$. Same models, same frozen sampler, same items — only T differs.

Table 1: *Fitted exponents on the 75-point surface* (M_{sep} for α, β, τ ; M_{sub} for κ) with 95% run-level bootstrap CIs for the step exponents. $\Delta\tau = 0.065$, CI [0.053, 0.077], excludes zero: the pre-registered primary result. κ is at its bound in both rows and $\Delta\text{AIC}_{\text{sub-sep}} \gg +4$, so no exchange rate is quotable. $\Delta\text{AIC}_{\text{full-N}} \leq -4$ means iso- N shape matters.

Metric	α (width)	β (depth)	τ (steps)	κ	$\Delta\text{AIC}_{\text{sub-sep}}$	$\Delta\text{AIC}_{\text{full-N}}$
WER	2.951	1.912	0.809 [0.803, 0.814]	3.000 [†]	+55.2	−29.5
SIM-o	0.917	0.672	0.744 [0.732, 0.755]	3.000 [†]	+89.8	−16.3

[†]at the fitted bound, i.e. not identified; $\Delta\tau = 0.065$, CI [0.053, 0.077].

for; WER does not extrapolate (24.7% versus 23.2%). Degenerate outputs — empty, truncated or looping transcripts under the frozen rule — fall monotonically with the step budget, from 41.0% at $T=1$ to 0.1% at $T=16$: most of what $T=1$ costs is not subtle quality but outright failures.

Serving corollary. Measured on the evaluation GPU at batch 1, c_{layer} is flat in width across the whole grid (0.493–0.514 ms per layer for $w \in [256, 1152]$): at these sizes the device is launch-bound, so serial latency is $8Td c_{\text{layer}}$ and width is free. The allocation rule follows: spend parameters on **depth** until latency binds, then spend the remaining latency on **steps**, which buy intelligibility to $T_{\text{WER}}^* = 16$ but stop buying identity after 8. A system that needs identity more than intelligibility can run at 8 steps and reclaim 50% of a $T=16$ configuration’s serial latency.

5 Limitations and outlook

The ceiling is 133 M non-embedding parameters and 30,000 steps; every model is under-trained by frontier standards, and the T curves may flatten differently at scale — $T_{\text{WER}}^* = 16$ is a lower bound, since WER had not saturated at the largest budget we measured. One language, one codec, one corpus, and no classifier-free guidance — guidance is a second inference-time axis that would confound the step axis, so excluding it is deliberate, and it is the obvious follow-up. The $\Delta\tau$ interval resamples 2 runs per config, so it captures seed-to-seed variation only, not the uncertainty of the single shared harness and corpus. $\Delta\rho$ and κ are reported as unidentified rather than estimated. Cuts taken: none. This is pre-registered outcome class S1, and selectivity is the only positive claim the data license.

Reproducibility. Protocol, hypotheses, thresholds, frozen sampler and analysis plan were registered before the first run; every gate is logged with its measured values. The released artefact holds the harness, all 30 run records, the 75-point surface, and the fitting and bootstrap code. Total cost 143 GPU-hours.

References

Philip Anastassiou et al. Seed-TTS: A family of high-quality versatile speech generation models. *arXiv preprint arXiv:2406.02430*, 2024.

146 Jacob Austin, Daniel D. Johnson, Jonathan Ho, Daniel Tarlow, and Rianne van den Berg. Structured
147 denoising diffusion models in discrete state-spaces. In *Advances in Neural Information Processing*
148 *Systems (NeurIPS)*, 2021.

149 Zalán Borsos, Matt Sharifi, Damien Vincent, Eugene Kharitonov, Neil Zeghidour, and Marco
150 Tagliasacchi. SoundStorm: Efficient parallel audio generation. *arXiv preprint arXiv:2305.09636*,
151 2023.

152 Huiwen Chang, Han Zhang, Lu Jiang, Ce Liu, and William T. Freeman. MaskGIT: Masked generative
153 image transformer. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*,
154 2022.

155 Sanyuan Chen, Chengyi Wang, Zhengyang Chen, Yu Wu, Shujie Liu, Zhuo Chen, Jinyu Li, Naoyuki
156 Kanda, Takuya Yoshioka, Xiong Xiao, Jian Wu, Long Zhou, Shuo Ren, Yanmin Qian, Yao Qian,
157 Michael Zeng, Xiangzhan Yu, and Furu Wei. WayLM: Large-scale self-supervised pre-training for
158 full stack speech processing. *IEEE Journal of Selected Topics in Signal Processing*, 2022.

159 Alexandre Défossez, Laurent Mazaré, Manu Orsini, Amélie Royer, Patrick Pérez, Hervé Jégou,
160 Edouard Grave, and Neil Zeghidour. Moshi: a speech-text foundation model for real-time dialogue.
161 *arXiv preprint arXiv:2410.00037*, 2024.

162 Jordan Hoffmann, Sebastian Borgeaud, Arthur Mensch, Elena Buchatskaya, Trevor Cai, Eliza
163 Rutherford, Diego de Las Casas, Lisa Anne Hendricks, Johannes Welbl, Aidan Clark, et al.
164 Training compute-optimal large language models. In *Advances in Neural Information Processing*
165 *Systems (NeurIPS)*, 2022.

166 Jared Kaplan, Sam McCandlish, Tom Henighan, Tom B. Brown, Benjamin Chess, Rewon Child,
167 Scott Gray, Alec Radford, Jeffrey Wu, and Dario Amodei. Scaling laws for neural language models.
168 *arXiv preprint arXiv:2001.08361*, 2020.

169 Shen Nie, Fengqi Zhu, Zebin You, Xiaolu Zhang, Jingyang Ou, Jun Hu, Jun Zhou, Yankai Lin, Ji-
170 Rong Wen, and Chongxuan Li. Large language diffusion models. *arXiv preprint arXiv:2502.09992*,
171 2025.

172 Alec Radford, Jong Wook Kim, Tao Xu, Greg Brockman, Christine McLeavey, and Ilya Sutskever.
173 Robust speech recognition via large-scale weak supervision. In *International Conference on*
174 *Machine Learning (ICML)*, 2023.

175 Takaaki Saeki, Detai Xin, Wataru Nakata, Tomoki Koriyama, Shinnosuke Takamichi, and Hiroshi
176 Saruwatari. UTMOS: UTokyo-SaruLab system for VoiceMOS challenge 2022. In *Interspeech*,
177 2022.

178 Subham Sekhar Sahoo, Marianne Arriola, Yair Schiff, Aaron Gokaslan, Edgar Marroquin, Justin T.
179 Chiu, Alexander Rush, and Volodymyr Kuleshov. Simple and effective masked diffusion language
180 models. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.

181 Charlie Snell, Jaehoon Lee, Kelvin Xu, and Aviral Kumar. Scaling LLM test-time compute optimally
182 can be more effective than scaling model parameters. *arXiv preprint arXiv:2408.03314*, 2024.

183 Chengyi Wang, Sanyuan Chen, Yu Wu, Ziqiang Zhang, Long Zhou, Shujie Liu, Zhuo Chen, Yanqing
184 Liu, Huaming Wang, Jinyu Li, Lei He, Sheng Zhao, and Furu Wei. Neural codec language models
185 are zero-shot text to speech synthesizers. *arXiv preprint arXiv:2301.02111*, 2023.

186 Yuancheng Wang, Haoyue Zhan, Liwei Liu, Ruihong Zeng, Haotian Guo, Jiachen Zheng, Qiang
187 Zhang, Xueyao Zhang, Shunsi Zhang, and Zhizheng Wu. MaskGCT: Zero-shot text-to-speech
188 with masked generative codec transformer. *arXiv preprint arXiv:2409.00750*, 2024.

189 Greg Yang, Edward J. Hu, Igor Babuschkin, Szymon Sidor, Xiaodong Liu, David Farhi, Nick
190 Ryder, Jakub Pachocki, Weizhu Chen, and Jianfeng Gao. Tensor programs V: Tuning large neural
191 networks via zero-shot hyperparameter transfer. In *Advances in Neural Information Processing*
192 *Systems (NeurIPS)*, 2021.

193 Zhen Ye, Xinfu Zhu, Chi-Min Chan, Xinsheng Wang, Xu Tan, Jiahe Lei, Yi Peng, Haohe Liu, Yizhu
194 Jin, Zheqi Dai, Hongzhan Lin, Jianyi Chen, Xingjian Du, Liumeng Xue, Yunlin Chen, Zhifei Li,
195 Lei Xie, Qiuqiang Kong, Yike Guo, and Wei Xue. Llasa: Scaling train-time and inference-time
196 compute for LLaMA-based speech synthesis. *arXiv preprint arXiv:2502.04128*, 2025.

197 A The shape grid, as run

198 Every config in the active grid with its measured $T=16$ means (seeds 0, 1, 400 eval_zs items each).
 199 Iso- N within a budget: the three budgets sit at 20, 50 and 125 M non-embedding parameters and the
 200 five shapes per budget span the aspect-ratio range at fixed $N = 12dw^2$.

config	budget	w	d	heads	N (M)	val loss	WER (%)	SIM-o	Degen (%)
A1	A	640	4	10	19.7	4.631	47.0	0.332	0.62
A2	A	448	8	7	19.3	4.570	22.1	0.351	0.12
A3	A	384	12	6	21.2	4.564	15.2	0.361	0.12
A4	A	320	18	5	22.1	4.578	14.4	0.355	0.00
A5	A	256	24	4	18.9	4.606	13.5	0.351	0.12
B1	B	832	6	13	49.8	4.548	26.5	0.362	0.00
B2	B	640	10	10	49.2	4.507	17.0	0.380	0.00
B3	B	512	16	8	50.3	4.498	12.8	0.389	0.12
B4	B	448	22	7	53.0	4.509	12.5	0.384	0.00
B5	B	384	30	6	53.1	4.515	10.0	0.377	0.00
C1	C	1152	8	18	127.4	4.494	15.7	0.379	0.25
C2	C	960	12	15	132.7	4.507	19.8	0.371	0.12
C3	C	768	18	12	127.4	4.484	15.4	0.397	0.00
C4	C	640	26	10	127.8	4.482	13.3	0.392	0.12
C5	C	512	36	8	113.2	4.486	10.1	0.385	0.00

202 B Additional figures




Figure 3: *Why the substitution form is rejected.* Iso-WER contours of M_{sub} in $(\log T, \log d)$ at the median width, with the iso-latency line $8Td = \text{const}$. The fitted κ sits on its bound and $\Delta\text{AICc}(M_{\text{sub}} - M_{\text{sep}}) = +55.2$, so these contours are shown as a rejected model, not as a design rule.

203 C Fit diagnostics and exploratory analyses

204 **Parameters at bounds.** The width coefficient A reaches its lower bound (0) in the $T=16$ Part-A fit
 205 for both metrics, and κ reaches its upper bound in every one of the 2000 bootstrap replicates of M_{sub} ;
 206 both are reported as non-identification rather than as estimates. The irreducible term E reaches 0 for
 207 WER under M_{sep} .

208 **Weighting.** Points are weighted $1/\text{SE}^2$ with SE taken over seeds, floored at a quarter of the pooled
 209 seed SD so that a chance agreement between two seeds cannot dominate the fit; the number of floored
 210 points is recorded in the released `fits.json`.

211 **Exploratory (not part of any declared claim).** An unweighted refit of both parts is released
 212 alongside the weighted fit, as is a fit that additionally includes the finished runs of a stretch third seed.
 213 Both reproduce the sign and significance of $\Delta\tau$.




Figure 4: *Fitted anisotropy surface* at $T=16$ for both metrics, with the measured config means overlaid. The width coefficient A collapses to its lower bound for both metrics, which is why $\rho = \alpha/\beta$ is reported as unidentified in §4 even though the two-axis model beats the N -only model ($\Delta\text{AICc} = -29.5$ and -16.3).




Figure 5: *H-D4 extrapolation*. The two smaller budgets are fitted and the largest budget predicted, under M_{full} and the N -only model. SIM-o extrapolates (MAPE 2.2% vs 3.0%); WER does not (24.7% vs 23.2%).

214 Use of Generative AI Disclosure

215 Generative AI was used as a coding assistant for the experimental harness and for language refinement
 216 of author-written text. The research design, the pre-registered hypotheses and thresholds, the analysis
 217 plan, and all scientific claims are the authors' own; the authors verified the code and the numbers and
 218 take full responsibility for this paper.