
Refinement Buys Intelligibility, Search Buys Identity: Per-Metric Test-Time Scaling for Masked-Diffusion TTS

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 A diffusion language model buys sequential computation twice: depth, paid once
2 in parameters, and denoising steps, rented per utterance at serve time. We ask
3 whether the rented kind helps every capability equally. We train an iso-parameter
4 width–depth grid of 15 masked-diffusion codec TTS models (3 seeds, 20–125 M
5 non-embedding parameters) and sweep the refinement budget T at inference only,
6 giving a 75-point (w, d, T) surface per metric from 45 trained models. Fitting
7 $\text{err} = E + Aw^{-\alpha} + Bd^{-\beta} + CT^{-\tau}$ pre-registered, test-time scaling is metric-
8 selective: $\Delta\tau = \tau_{\text{WER}} - \tau_{\text{SIM}} = 0.110$, 95% run-level bootstrap CI [0.092, 0.131],
9 excluding zero. The effect that carries practical weight is one of *magnitude*: over
10 $T=1 \rightarrow 16$ refinement moves 0.952 absolute WER (a $6.1\times$ reduction) against
11 0.161 absolute speaker-similarity error ($1.25\times$), a $5.9\times$ asymmetry. It is not a
12 difference in convergence *rate*: on each metric’s own normalised curve the fraction
13 of total gain realised by $T=8$ is 94.6% for WER and 93.7% for identity, identity
14 marginally later. The selectivity persists across a 276 M four-budget replication,
15 $\Delta\tau = +0.1095$, CI [0.0951, 0.1241]. The asymmetry is less a limit on what
16 inference can buy than on what *refinement* can buy: at matched NFE, best-of-
17 K search moves speaker similarity by +0.0444 (CI [0.0416, 0.0470]) against
18 refinement at $T=64$, or +0.0365 against the deployable $T=16$ default, where
19 refinement moves it least. Depth and steps are *not* interchangeable: the substitution
20 form $B(dT^\kappa)^{-\beta}$ is worse than the separable form by $\Delta\text{AICc} = +69.3$.

21 1 Introduction

22 A diffusion language model can buy sequential computation twice. It *owns* depth: d transformer
23 blocks, paid once in parameters and then on every forward pass forever. It *rents* refinement steps:
24 T denoising passes over the same sequence, paid per utterance and adjustable at serve time. An
25 autoregressive model has only the first lever; its serial cost per token is fixed by the architecture the
26 moment training ends. If rented serial compute substitutes for owned serial compute, a deployment
27 can trade parameters for steps after the fact — but only for the capabilities that steps actually improve.

28 Zero-shot text-to-speech is a clean probe: intelligibility is compositional (the right phonetic content in
29 the right order, read out by word error rate) while speaker identity is representational (timbre carried
30 from a short prompt, read out by SIM-o). The two need not respond to the same knob.

31 We train one iso-parameter width–depth grid of masked-diffusion codec language models and then
32 sweep T at inference only, which yields a full (w, d, T) response surface per metric at zero additional
33 training cost. Two claims were pre-registered before any run. **H-D2**: refinement steps improve
34 intelligibility more than identity, i.e. test-time scaling is metric-selective. **H-D3**: for intelligibility,
35 steps and depth exchange at a measurable rate κ — how many step-doublings buy one depth-doubling.




Figure 1: *The two ways a masked-diffusion TTS model buys sequential computation.* Phonemes and a 3-second prompt are always visible; the generated region of the $8 \times \text{frames}$ token grid starts fully **masked**. The backbone spends **depth** — *owned* serial compute, fixed at training time — and the sampler spends **T refinement steps per codebook level** — *rented* serial compute, chosen at serve time. One training grid over (w, d) therefore yields a whole (w, d, T) surface at inference. This paper asks whether the rented kind helps every capability equally.

$$\begin{aligned} \text{separable: } \text{err}_m(w, d, T) &= E_m + A_m w^{-\alpha_m} + B_m d^{-\beta_m} + C_m T^{-\tau_m} \\ \text{substitution: } \text{err}_m(w, d, T) &= E_m + A_m w^{-\alpha_m} + B_m (d T^{\kappa_m})^{-\beta_m} \end{aligned} \quad (1)$$

36 The first answer is yes and quantitative: $\Delta\tau = 0.110$, run-level bootstrap CI $[0.092, 0.131]$. The
37 second is a clean no — the substitution form is decisively *worse* than the separable one ($\Delta\text{AICc} =$
38 $+69.3$ for WER, $+71.3$ for SIM-o) and κ runs to its bound, so no finite depth–step exchange
39 rate describes this regime. Extra inference compute is a targeted instrument, not a substitute for
40 architecture.

41 ❶ The first per-metric (w, d, T) scaling surface for speech generation: 75 points per metric from 45
42 trained models, released with the harness. ❷ $\Delta\tau = 0.110$, CI $[0.092, 0.131]$ — refinement steps are
43 metric-selective — replicated across four parameter budgets, $3\times$ training compute, a second ASR
44 and an unbounded reparameterisation, with depth–step substitution reported as a *negative* AICc result
45 and κ as unidentified. ❸ An identity ledger measured against a codec ceiling, showing what does
46 move speaker similarity when refinement does not: training compute and test-time selection, with
47 prompt context and rate matching buying nothing.

48 2 Related work

49 Masked diffusion language models are trained at scale [Nie et al., 2025, Sahoo et al., 2024, Austin
50 et al., 2021], but their scaling is reported against a single likelihood-shaped objective, so a per-
51 capability question cannot arise. The masked-generative recipes that made codec-token TTS practical
52 [Borsos et al., 2023, Chang et al., 2022, Wang et al., 2024] fixed a step budget by engineering
53 judgement and report a single operating point; none maps error against steps as a law. Autoregressive
54 codec TTS [Wang et al., 2023] reports scaling in total parameters only [Ye et al., 2025], which
55 separates neither width from depth nor either from steps, as is true of classical scaling work [Kaplan
56 et al., 2020, Hoffmann et al., 2022]. Test-time scaling is well studied in reasoning [Snell et al.,
57 2024], where the payoff is one correctness metric; whether extra inference compute helps different
58 capabilities *differently* is, to our knowledge, untested in generative speech.

59 3 Method

60 **Backbone and objective.** A bidirectional (non-causal) pre-LN transformer over Mimi [Défossez
61 et al., 2024] codec tokens at 12.5 Hz with the first 8 residual codebooks (2048 entries each, plus
62 MASK and PAD). **Width** w , **depth** d , head dimension 64, GELU MLP ratio 4, RMSNorm, RoPE,
63 no weight tying, so the non-embedding parameter count is $N = 12 d w^2$. One embedding table
64 per codebook is summed at the input and one linear head per codebook reads out. Training is
65 SoundStorm-style [Borsos et al., 2023] coarse-to-fine masking: per example a level $l \sim \mathcal{U}\{1..8\}$
66 and a ratio $t \sim \mathcal{U}(0, 1]$; levels below l stay visible, level l is masked i.i.d. with probability t , levels
67 above l are entirely MASK; the loss is cross-entropy on the masked cells of level l only. A phoneme
68 prefix and a 3-second prompt of codec tokens are always visible; there is no speaker encoder, so the
69 backbone is the only identity pathway, and no classifier-free guidance, which would confound the

Table 1: $\Delta\tau$ under every variation we could apply. The headline holds across a fourth parameter budget, 3 $\times$ training compute, a different ASR, and an unbounded reparameterisation of the fitted amplitudes.

variation	what changes	$\Delta\tau$	95% CI
none (75 points, 45 runs)	—	0.110	[0.092, 0.131]
fourth budget, 276 M	10 further runs	+0.1095	[0.0951, 0.1241]
3 $\times$ training compute	3 configs at 90k steps	+0.1767	[0.0789, 0.2854]
ASR swapped	Whisper-medium.en	+0.1283	[0.1064, 0.1492]
unbounded amplitudes	log A , log B , log C free	+0.1260	[0.1043, 0.1474]

step axis. The target length is fixed before generation from the text length and a corpus character rate (0.0602 s/char).

Frozen sampler. MaskGIT-style confidence decoding [Chang et al., 2022], level by level, with exactly T refinement steps per level, so $\text{NFE} = 8T$: a cosine schedule sets how many cells remain masked after each step, tokens are drawn at temperature 1.0, and confidences are log-probabilities plus Gumbel noise annealed $1 \rightarrow 0$; committed cells stay committed. The sampler is identical for every config, seed and T , and generation noise is keyed so that T is the only difference between surface points of one run. As a sanity check, $T=1$ and $T=16$ outputs differ on more than 20% of generated token cells, confirming that T reaches the sampler.

Design and statistics. 15 configs span three iso- N budgets $\times$ five aspect ratios, each trained at 3 seeds for 30,000 steps at effective batch 256 with μP [Yang et al., 2021] (base width 256, base LR 0.004; transfer verified by five-point LR sweeps). $T \in \{1, 2, 4, 8, 16\}$ is swept at evaluation only, giving 75 surface points per metric from 45 models. Metrics: WER (Whisper-large-v3 [Radford et al., 2023], greedy), SIM-o (WavLM-large [Chen et al., 2022] against the *original* prompt waveform, following Anastassiou et al. [2024]), a frozen degenerate-output rule, and UTMOS [Saeki et al., 2022]. Error transforms are $\text{err}_{\text{WER}} = \text{WER}$ and $\text{err}_{\text{SIM}} = 1 - \text{SIM-o}$, computed over all items including degenerates. Both forms of Eq. (1) are fitted by weighted non-linear least squares with 32 multi-starts; $\Delta\tau = \tau_{\text{WER}} - \tau_{\text{SIM}}$, κ , and $\Delta\rho = \alpha_{\text{SIM}}/\beta_{\text{SIM}} - \alpha_{\text{WER}}/\beta_{\text{WER}}$ carry percentile confidence intervals from a bootstrap that resamples at the *run* level — the five T -points of one run share a trained model, so resampling surface points would understate their correlation — with 2000 replicates recomputing every T point.

4 Results

Steps are metric-selective (H-D2, supported). Fitted on all 75 surface points, the step exponent is $\tau_{\text{WER}} = 0.836$ (CI [0.825, 0.851]) against $\tau_{\text{SIM}} = 0.726$ (CI [0.709, 0.744]), so $\Delta\tau = 0.110$ with CI [0.092, 0.131], excluding zero. The raw surface shows it without any model: over the grid, $T=1 \rightarrow T=16$ takes WER from 113.8% to 18.7% (a factor of $6.1\times$ on the error scale) while SIM-o moves only $0.209 \rightarrow 0.370$ ($1.25\times$ on $1 - \text{SIM-o}$), an absolute-error asymmetry of $5.9\times$. The selectivity is not an artefact of scale, training compute, or the metric stack: a fourth budget (276 M, ten further runs) gives $\Delta\tau = +0.1095$, CI [0.0951, 0.1241]; retraining three configurations $3\times$ longer gives $+0.1767$, CI [0.0789, 0.2854]; and a different ASR or unbounded log-amplitudes move it to $+0.1283$ and $+0.1260$. Extending T to 64 shows WER still falling past the grid edge ($0.1516 \rightarrow 0.1294$), so the $T=16$ value was a lower bound. The power gate passed in all three budgets before these tests ran.

Depth and steps do not substitute (H-D3, not supported). The pre-registered substitution form collapses: $\Delta\text{AICc}(M_{\text{sub}} - M_{\text{sep}}) = +69.3$ for WER and $+71.3$ for SIM-o where the hypothesis required $\leq +4$, and κ sits on its bound for both metrics, leaving the exchange rate unidentified. The fits show why: T enters as an additive term decaying to a floor, while depth moves the floor itself, and one factor dT^κ cannot fit both orderings.

Shape matters, but the anisotropy ratio is not identified (H-D1, not supported). At $T=16$ the two-axis model beats the N -only model for both metrics ($\Delta\text{AICc} = -15.5 / -5.5$ against a -4 threshold), so error is not a function of parameter count alone. The *direction* of the anisotropy we




Figure 2: *Test-time scaling is metric-selective.* Error versus **refinement steps** for all 15 configs, each normalised to its own $T=1$ value; thick lines are grid means. **(A)** WER error falls 0.9518 in absolute terms over $T=1 \rightarrow 16$ and is still falling at $T=16$. **(B)** SIM-o error falls 0.1606 over the same range, $5.9\times$ less. Same models, same frozen sampler, same items — only T differs.

Table 2: *Fitted exponents* on the 75-point surface (M_{sep} for α, β, τ ; M_{sub} for κ) with 95% run-level bootstrap CIs for the step exponents. $\Delta\tau = 0.110$, CI [0.092, 0.131], excludes zero: the pre-registered primary result. κ is at its bound in both rows and $\Delta\text{AICc}_{\text{sub-sep}} \gg +4$, so no exchange rate is quotable. $\Delta\text{AICc}_{\text{full-N}} \leq -4$ means iso- N shape matters.

Metric	α (width)	β (depth)	τ (steps)	κ	$\Delta\text{AICc}_{\text{sub-sep}}$	$\Delta\text{AICc}_{\text{full-N}}$
WER	1.574	2.040	0.836 [0.825, 0.851]	3.000 [†]	+69.3	−15.5
SIM-o	0.919	0.553	0.726 [0.709, 0.744]	3.000 [†]	+71.3	−5.5

[†]at the fitted bound, i.e. not identified; $\Delta\tau = 0.110$, CI [0.092, 0.131].

cannot claim: $\Delta\rho = 0.174$, CI [−0.711, 0.388] spanning zero, because the width amplitude saturates its pre-registered bound in every fit, leaving only $A w^{-\alpha}$ identified. **Depth** is identified and separates the capabilities sharply ($\beta_{\text{WER}} = 2.040$ versus $\beta_{\text{SIM}} = 0.553$): moving from the shallowest shape to the best depth is worth 40.2, 16.3 and 7.3 WER points across the three budgets, while SIM-o varies by under 0.065. The depth optimum is interior only in the smallest budget; in the two larger ones the argmin sits at the deepest shape tested, so these are lower bounds rather than interior optima.

Extrapolation and reliability. Fitting the two smaller budgets to predict the largest gives MAPE 2.4% for SIM-o (versus 2.2% N -only), inside the 15% asked for; WER does not extrapolate (20.2% versus 20.4%). Degeneracy falls monotonically with T , 40.4% to 0.1%: most of what $T=1$ costs is outright failure.

Serving corollary. Measured at batch 1, c_{layer} is flat in width (0.493–0.514 ms per layer for $w \in [256, 1152]$): the device is launch-bound, so serial latency is $8Td c_{\text{layer}}$ and width is free, leaving a latency budget to constrain only the product Td . Our data do not resolve how to split it — the pre-registered depth-before-steps rule held at 5.9% of 40 budgets against an 80% bar, and T is censored at 16, so the data do not favour either direction.

What buys identity, measured against a ceiling. If refinement is an intelligibility knob, what moves identity? We measure the candidates on one scale against a bound: a codec round-trip of ground-truth audio scores 0.555, leaving 0.0747 of headroom above our best system. The largest lever is *training compute* (30k→90k steps, +0.0827, CI [0.0733, 0.0908] over 9 runs; a further doubling adds +0.0155). Next is *test-time search*: at matched NFE ($8T$ per candidate), best-of- K sampling with candidates chosen by one speaker encoder and scored by a second, independently gated one (§ Limitations) improves similarity by +0.0444 (CI [0.0416, 0.0470], 72 % of items) over refinement at $T=64$, or +0.0365 over the deployable $T=16$ default for +0.0425 WER — while refinement stays the better spend for intelligibility (WER 0.1316 versus 0.1955). The gain is selection, not sampling: an arbitrary candidate scores 0.4259 regardless of K . It persists at 276 M non-embedding parameters (+0.0422) and, on three runs, at $3\times$ training compute (+0.0368). Further behind sit a $14\times$ parameter increase (+0.0531) and training-free speaker-contrastive guidance (+0.0121, for +2.73 WER points — a trade), while rate-matched length conditioning (+0.0044) and longer prompts buy nothing: the best prompt is the 3 s default seen in training, and beyond it output stays fluent and

Table 3: *What buys speaker identity.* Absolute SIM-o gain per axis, measured against the 0.0747 of headroom that remains between our best system and a codec round-trip of real audio. Training compute and test-time selection dominate; two test-time levers buy nothing.

axis	change	Δ SIM-o	cost
training compute	30k $\rightarrow$ 90k steps	+0.0827	3 $\times$ training
(further)	90k $\rightarrow$ 180k	+0.0155	2 $\times$ again
test-time search	best-of-8 vs $T=16$	+0.0365	+0.0425 WER, 8 $\times$ NFE
parameters	20 M $\rightarrow$ 276 M	+0.0531	14 $\times$ parameters
guidance	$\gamma = 0.5$	+0.0121	+2.73 WER points
rate matching	per-item duration	+0.0044	free
prompt context	3 s $\rightarrow$ 9 s	none	+73 WER points




Figure 3: *Identity is mostly not the model's to win.* (A) what each axis buys in speaker similarity, against the headroom that remains to the codec ceiling. (B) of the distance between real audio and our best system, 62% is lost in the codec itself.

140 speaker-consistent while decoupling from the text. Identity is bought with training and selection, not
 141 with more refinement.

142 Where the passes are spent matters as much as how many: at matched NFE = 32, moving them to the
 143 coarsest codebook level changes WER from 1.1480 to 0.2191, at a cost in identity and naturalness
 144 (Appendix D).

145 5 Limitations and outlook

146 Our ceiling is 133 M non-embedding parameters and 30,000 steps, so every model is under-trained
 147 and T_{WER}^* is a lower bound. Part of the identity gap is not ours to close: a codec round-trip of the
 148 target scores 0.555 against 0.678 for the original, so 62% of the distance to real speech is lost in
 149 the codec. One limitation of the search result remains open: selector and scorer are both ECAPA-
 150 TDNN-family encoders (item-level $r = 0.71 - -0.73$); each passes the ground-truth speaker gate
 151 independently and the selector captures only 46 % of an oracle's gap where shared bias would predict
 152 nearly all of it, but excluding representational overlap needs a different scorer family, and matching is
 153 on diffusion NFE only. One language, one codec, one corpus, and no classifier-free guidance. The
 154 $\Delta\tau$ interval captures seed variation only; target length comes from a corpus median; $\Delta\rho$ and κ are
 155 unidentified. Cuts taken: none; pre-registered outcome class S1.

156 References

- 157 Philip Anastassiou et al. Seed-TTS: A family of high-quality versatile speech generation models.
 158 *arXiv preprint arXiv:2406.02430*, 2024.
- 159 Jacob Austin, Daniel D. Johnson, Jonathan Ho, Daniel Tarlow, and Rianne van den Berg. Structured
 160 denoising diffusion models in discrete state-spaces. In *Advances in Neural Information Processing*

Table 4: *Refinement and search buy different capabilities at the same cost.* Matched NFE (8T per candidate); best-of- K candidates selected by WavLM-SV and scored by ECAPA, which passes the ground-truth speaker gate independently. Refinement wins intelligibility at every tier; search wins identity at every tier.

NFE	refinement	search	Δ ECAPA	95% CI	WER (refine / search)
128	$T=16$	$K=2$	+0.0153	[0.0127, 0.0179]	0.1530 / 0.1938
256	$T=32$	$K=4$	+0.0322	[0.0307, 0.0338]	0.1335 / 0.1947
512	$T=64$	$K=8$	+0.0444	[0.0416, 0.0470]	0.1316 / 0.1955
at NFE 512, replicated out of scope: budget D +0.0422 [0.0393, 0.0448]; $3 \times$ training +0.0368 [0.0321, 0.0428]					

Table 5: *Where the refinement passes are spent matters as much as how many.* Three frozen per-level schedules at matched total NFE = 32 on the 125 M budget, 15 runs, 400 items. Spending every pass on the coarsest codebook level rather than the finest cuts WER $5.2 \times$; it is not free, costing identity and naturalness against a uniform split.

schedule	WER	SIM-o	UTMOS	degenerate
coarse [25, 1, . . . , 1]	0.2191	0.3057	2.17	0.1%
uniform [4, . . . , 4]	0.3144	0.3477	2.51	0.5%
fine [1, . . . , 1, 25]	1.1480	0.2246	1.56	39.1%

161 *Systems (NeurIPS)*, 2021.

162 Zalán Borsos, Matt Sharifi, Damien Vincent, Eugene Kharitonov, Neil Zeghidour, and Marco
163 Tagliasacchi. SoundStorm: Efficient parallel audio generation. *arXiv preprint arXiv:2305.09636*,
164 2023.

165 Huiwen Chang, Han Zhang, Lu Jiang, Ce Liu, and William T. Freeman. MaskGIT: Masked generative
166 image transformer. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*,
167 2022.

168 Sanyuan Chen, Chengyi Wang, Zhengyang Chen, Yu Wu, Shujie Liu, Zhuo Chen, Jinyu Li, Naoyuki
169 Kanda, Takuya Yoshioka, Xiong Xiao, Jian Wu, Long Zhou, Shuo Ren, Yanmin Qian, Yao Qian,
170 Michael Zeng, Xiangzhan Yu, and Furu Wei. WavLM: Large-scale self-supervised pre-training for
171 full stack speech processing. *IEEE Journal of Selected Topics in Signal Processing*, 2022.

172 Alexandre Défossez, Laurent Mazaré, Manu Orsini, Amélie Royer, Patrick Pérez, Hervé Jégou,
173 Edouard Grave, and Neil Zeghidour. Moshi: a speech-text foundation model for real-time dialogue.
174 *arXiv preprint arXiv:2410.00037*, 2024.

175 Jordan Hoffmann, Sebastian Borgeaud, Arthur Mensch, Elena Buchatskaya, Trevor Cai, Eliza
176 Rutherford, Diego de Las Casas, Lisa Anne Hendricks, Johannes Welbl, Aidan Clark, et al.
177 Training compute-optimal large language models. In *Advances in Neural Information Processing*
178 *Systems (NeurIPS)*, 2022.

179 Jared Kaplan, Sam McCandlish, Tom Henighan, Tom B. Brown, Benjamin Chess, Rewon Child,
180 Scott Gray, Alec Radford, Jeffrey Wu, and Dario Amodei. Scaling laws for neural language models.
181 *arXiv preprint arXiv:2001.08361*, 2020.

182 Shen Nie, Fengqi Zhu, Zebin You, Xiaolu Zhang, Jingyang Ou, Jun Hu, Jun Zhou, Yankai Lin, Ji-
183 Rong Wen, and Chongxuan Li. Large language diffusion models. *arXiv preprint arXiv:2502.09992*,
184 2025.

185 Alec Radford, Jong Wook Kim, Tao Xu, Greg Brockman, Christine McLeavey, and Ilya Sutskever.
186 Robust speech recognition via large-scale weak supervision. In *International Conference on*
187 *Machine Learning (ICML)*, 2023.

188 Takaaki Saeki, Detai Xin, Wataru Nakata, Tomoki Koriyama, Shinnosuke Takamichi, and Hiroshi
189 Saruwatari. UTMOS: UTokyo-SaruLab system for VoiceMOS challenge 2022. In *Interspeech*,
190 2022.

Table 6: *Longer prompts do not buy identity, and destroy intelligibility.* Prompt length swept at inference on 15 125 M-budget runs, 397 items, every arm scored against the same fixed 3 s reference. Training used 3 s prompts only; the 6 and 9 s arms test extrapolation, where output remains fluent and speaker-consistent while decoupling from the conditioning text.

prompt	SIM-o	WER
1.5 s	0.3213	0.1578
3 s (<i>training length</i>)	0.3839	0.1474
6 s	0.3724	0.7692
9 s	0.3759	0.8642

Subham Sekhar Sahoo, Marianne Arriola, Yair Schiff, Aaron Gokaslan, Edgar Marroquin, Justin T. Chiu, Alexander Rush, and Volodymyr Kuleshov. Simple and effective masked diffusion language models. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.

Charlie Snell, Jaehoon Lee, Kelvin Xu, and Aviral Kumar. Scaling LLM test-time compute optimally can be more effective than scaling model parameters. *arXiv preprint arXiv:2408.03314*, 2024.

Chengyi Wang, Sanyuan Chen, Yu Wu, Ziqiang Zhang, Long Zhou, Shujie Liu, Zhuo Chen, Yanqing Liu, Huaming Wang, Jinyu Li, Lei He, Sheng Zhao, and Furu Wei. Neural codec language models are zero-shot text to speech synthesizers. *arXiv preprint arXiv:2301.02111*, 2023.

Yuancheng Wang, Haoyue Zhan, Liwei Liu, Ruihong Zeng, Haotian Guo, Jiachen Zheng, Qiang Zhang, Xueyao Zhang, Shunsi Zhang, and Zhizheng Wu. MaskGCT: Zero-shot text-to-speech with masked generative codec transformer. *arXiv preprint arXiv:2409.00750*, 2024.

Greg Yang, Edward J. Hu, Igor Babuschkin, Szymon Sidor, Xiaodong Liu, David Farhi, Nick Ryder, Jakub Pachocki, Weizhu Chen, and Jianfeng Gao. Tensor programs V: Tuning large neural networks via zero-shot hyperparameter transfer. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2021.

Zhen Ye, Xinfu Zhu, Chi-Min Chan, Xinsheng Wang, Xu Tan, Jiahe Lei, Yi Peng, Haohe Liu, Yizhu Jin, Zheqi Dai, Hongzhan Lin, Jianyi Chen, Xingjian Du, Liumeng Xue, Yunlin Chen, Zhifei Li, Lei Xie, Qiuqiang Kong, Yike Guo, and Wei Xue. Llasa: Scaling train-time and inference-time compute for LLaMA-based speech synthesis. *arXiv preprint arXiv:2502.04128*, 2025.

A Reproducibility

Protocol, hypotheses, thresholds, the frozen sampler and the analysis plan were registered before the first run, and every gate is logged with its measured values. The released artefact holds the harness, all 45 run records, the 75-point surface, and the fitting and bootstrap code. Total cost 143 GPU-hours.

B The shape grid, as run

Every config in the active grid with its measured $T=16$ means (seeds 0, 1, 2, 400 eval_zs items each). Iso- N within a budget: the three budgets sit at 20, 50 and 125 M non-embedding parameters and the five shapes per budget span the aspect-ratio range at fixed $N = 12dw^2$.

218

config	budget	w	d	heads	N (M)	val loss	WER (%)	SIM-o	Degen (%)
A1	A	640	4	10	19.7	4.638	53.9	0.327	0.67
A2	A	448	8	7	19.3	4.582	28.4	0.348	0.08
A3	A	384	12	6	21.2	4.566	15.9	0.360	0.17
A4	A	320	18	5	22.1	4.574	13.6	0.359	0.00
A5	A	256	24	4	18.9	4.611	15.3	0.351	0.17
B1	B	832	6	13	49.8	4.551	27.8	0.361	0.08
B2	B	640	10	10	49.2	4.505	16.0	0.381	0.00
B3	B	512	16	8	50.3	4.497	11.9	0.388	0.08
B4	B	448	22	7	53.0	4.511	12.0	0.382	0.08
B5	B	384	30	6	53.1	4.520	11.5	0.375	0.08
C1	C	1152	8	18	127.4	4.504	17.1	0.373	0.17
C2	C	960	12	15	132.7	4.504	18.9	0.373	0.17
C3	C	768	18	12	127.4	4.488	14.3	0.392	0.00
C4	C	640	26	10	127.8	4.489	13.4	0.390	0.08
C5	C	512	36	8	113.2	4.487	9.9	0.388	0.00

219 C Additional figures




Figure 4: *Why the substitution form is rejected.* Iso-WER contours of M_{sub} in $(\log T, \log d)$ at the median width, with the iso-latency line $8Td = \text{const}$. The fitted κ sits on its bound and $\Delta\text{AICc}(M_{\text{sub}} - M_{\text{sep}}) = +69.3$, so these contours are shown as a rejected model, not as a design rule.




Figure 5: *Fitted anisotropy surface at $T=16$ for both metrics, with the measured config means overlaid.* The width amplitude A saturates its pre-registered upper bound for both metrics, which is why $\rho = \alpha/\beta$ is reported as unidentified in §4 even though the two-axis model beats the N -only model ($\Delta\text{AICc} = -15.5$ and -5.5).




Figure 6: *H-D4 extrapolation*. The two smaller budgets are fitted and the largest budget predicted, under M_{full} and the N -only model. SIM-o extrapolates (MAPE 2.4 % vs 2.2 %); WER does not (20.2 % vs 20.4 %).

D Refinement-budget allocation across codebook levels

Three frozen per-level step schedules at matched total NFE = 32 on the 125 M budget (15 runs, 400 items). Paired item-level bootstrap, 2000 replicates, pairing on the eval item so item difficulty cancels. Coarse spends all passes on codebook level 0, fine on level 7, uniform splits evenly.

schedule	WER	SIM-o	UTMOS	degenerate
coarse [25, 1, . . . , 1]	0.2191	0.3057	2.17	0.1%
uniform [4, . . . , 4]	0.3144	0.3477	2.51	0.5%
fine [1, . . . , 1, 25]	1.1480	0.2246	1.56	39.1%

Allocation is metric-selective in the same way the step *count* is: coarse wins WER but loses identity and naturalness against uniform, both with intervals excluding zero. It is a trade, not a free win.

A significance-based saturation contrast. Defining saturation as the smallest T whose paired-bootstrap interval for $\text{err}(T) - \text{err}(64)$ covers zero gives $T_{\text{WER}}^* = 32$ against $T_{\text{SIM}}^* = 16$ on the extended grid. We report this only here and do not build on it: roughly half the contrast is a precision artefact, since identity’s detection floor is $\approx 1.8\times$ looser relative to its own dynamic range than intelligibility’s. The raw-scale T^* — 95% of the $T=16$ value — is not affine-invariant and must never be compared across metrics: the same 5% band is 0.67% of WER’s range but 11.96% of SIM-o’s. On the invariant fraction-of-total-gain definition, $T^* = 16$ for both metrics.

E Fit diagnostics and exploratory analyses

Parameters at bounds. The width amplitude A saturates the pre-registered *upper* bound $A \leq 10$ in every fit (Part A: 10.000 for both metrics; Part B: 9.999 for WER, 10.000 for SIM-o), so only the product $A w^{-\alpha}$ is identified and α is not; the bound, not the data, is binding, which is a limitation of the pre-registered parameterisation worth fixing in any follow-up. κ reaches its upper bound in every one of the 2000 bootstrap replicates of M_{sub} . The irreducible term E reaches 0 for WER under M_{sep} . The depth amplitude B is interior throughout, so β is identified.

Exploratory: the depth exponent separates the capabilities. $\beta_{\text{WER}} = 2.040$ against $\beta_{\text{SIM}} = 0.553$, a factor of 3.7 — the asymmetry H-D1 sought, in the parameter the data pin down. Part-B ρ values point the pre-registered way ($\Delta\rho = +0.890$) but inherit the A -cap problem; the pre-registered test is Part A, so this is exploratory only.

Weighting. Points are weighted $1/\text{SE}^2$ with SE taken over seeds, floored at a quarter of the pooled seed SD so that a chance agreement between two seeds cannot dominate the fit; the number of floored points is recorded in the released `fits.json`.

248 **Exploratory (not part of any declared claim).** An unweighted refit of both parts is released
249 alongside the weighted fit, as is a fit that additionally includes the finished runs of a stretch third seed.
250 Both reproduce the sign and significance of $\Delta\tau$.

251 **Use of Generative AI Disclosure**

252 Generative AI was used as a coding assistant for the experimental harness and for language refinement
253 of author-written text. The research design, the pre-registered hypotheses and thresholds, the analysis
254 plan, and all scientific claims are the authors' own; the authors verified the code and the numbers and
255 take full responsibility for this paper.